

License Plate Detection and Text Recognition

Feb 17 – Jun 4 2020

AT THE VERY EDGE

Enabling AI

Presented by:

Lorenzo Lamberti
GreenWaves Technologies

Project Overview

Project Goal:

Application of Optical Character Recognition (OCR) techniques to **implement license plate recognition on GAP.**



京QC8152

Project Flow:

1. Study of related works:

find a suitable algorithm for embedded devices

2. Implementation of the algorithm (python code)

3. Evaluation of performances:

accuracy and reliability of model's predictions

4. Optimization of the model to satisfy GAP's constraints

Introduction to OCR

Optical Character Recognition

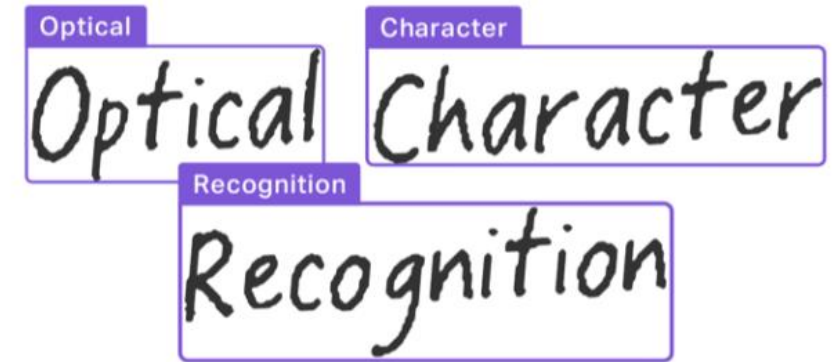
Goal: recognize text in images and transcribe it .

OCR categories:

- **In the wild** (unpredictable scenario)
- **controlled environment** (fixed setup)

License plate recognition is a sub-category of OCR in the wild

OCR in the wild is the hardest task and not completely solved



Challenges OCR in the wild

- **Artifacts:** light variations and reflections.
- **Weather conditions:** snow and rain.
- **Text is sparse:** text can be randomly located in the image.
- **Font size not constant:** depends on the object distance.
- **Variability of fonts** (including special characters e.g. logograms for China)




OCR Strategies

Traditional approaches:

Classic computer vision techniques (Tesseract)
(preprocessing + template/pattern matching)

- Based on a-priori assumptions
- Poor generalization on new scenarios: **bad results in the wild**

Input sample character



B mask

1	1	1	1	0
1	0	0	0	1
1	1	1	1	0
1	0	0	0	1
1	0	0	0	1
1	1	1	1	0

Result


1	1	1	1	
1				1
1	1	1	1	
1				1
1				1
1	1	1	1	

B

New approaches:

Based on **Deep learning** (Convolutional Neural Networks)

- Algorithm automatically learns features from given instances
- Good generalization if a large dataset is provided



State of the art

License plate Recognition

Standard Deep Learning approach to License Plate OCR:

The problem is tackled in **2 steps**:
(They are treated separately)


1. Text Detection:

Localization of the full words



2. Text Recognition:

Reconstruction of the word
character by character



Step 1: License Plate Detection


Object Detection algorithms:
(based on CNNs, convolutional neural networks)

- Single Shot Detector (**SSD**)
- You Only Look Once (YOLO)
- Region Based CNNs (R-CNN)


Classification VS Object Detection



CAT



CAT, DOG, DUCK



Step 2: License Plate Recognition



Recurrent Neural Networks (RNNs):

- LSTM
- BiLSTM
- CRNN

It is a type of network in which inside loops are allowed: memory effects.

It models temporal/spatial sequences.

It's the most common approach for text recognition
but Gap doesn't support RNNs at the moment



Step 2: License Plate Recognition

Convolutional Neural Networks Only:

- **Character segmentation with object detectors:**


Too redundant.

Need training data with segmented characters annotation: very rare

Character Segmentation



ABC1234



- **Holistic CNNs:**

8 branches of the CNN search for different positions in the image.

One branch specializes for detecting different characters in the plate.

Code not publicly available




- **LPRNet:**

End-to-end method without preliminary character segmentation. Straight from pixels to char predictions



<Beijing>FA9152



<Beijing>GZ5871

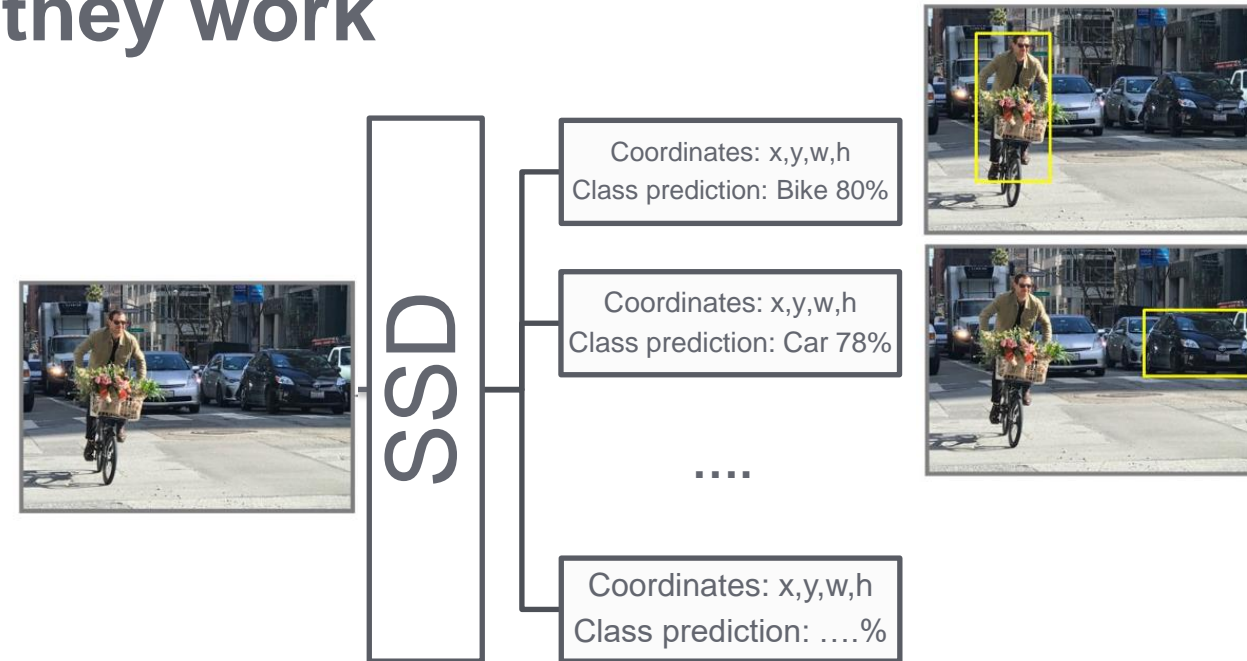


<Sichuan>A57237

Approaches selected: how they work

1) SSD: Single-Shot detector


- A fully Convolutional Neural Network
- Single step: simultaneous bounding box regression and classification
- One bbox and class associated to each object



2) LPRNet for character recognition:

Convolutional Neural network + Decoder

1. CNN predicts probability of characters and blank spaces
2. Decoder: associates probabilities to the right characters and removes double predictions



Approaches selected: **Why**

1° Step: **SSD** (object detection)

Code available in official Tensorflow obj detection API

Lightweight models available (Mobilenet)

Flexible: easy to switch architecture & a lot of pretrained models to try.

Big dataset ready to use.

License Plate Detection



2° Step: **LPRNet** (character recognition)

Code available (Intel - tensorflow)

Extremely lightweight: 1.7 Mb

Fast: inference time 3ms on a Nvidia GTX 1080

Promises an accuracy of over 95%

Only a CNN: doesn't use RNNs (not supported by gap)

Character Recognition



<Henan>K0F755

GAP Constraints:

- **Tensorflow** framework: we must provide a .tflite model
- **Memory: 8Mb** (must fit in L3 memory)
- **Real-Time Execution: ~ 1 FPS**
- Good accuracy (of course)

General design flow for ML algorithms

1. **Dataset search:** a lot data to train the algorithm
2. **Model selection:** satisfy GAP's constrains
3. Optimization: **quantization aware training**
4. **Validation:** test performances of real data

I acquired a little validation dataset (70 images) with GAP Himax camera

Validation dataset: GAP's Himax camera

Low resolution grayscale images (QVGA: 320x240)

20 Italian license plates



50 Chinese license plates



Acquired at 5 distances (0.3m, 1m, 2m, 3m, 4m)

0.3 m



1 m



2 m



3 m

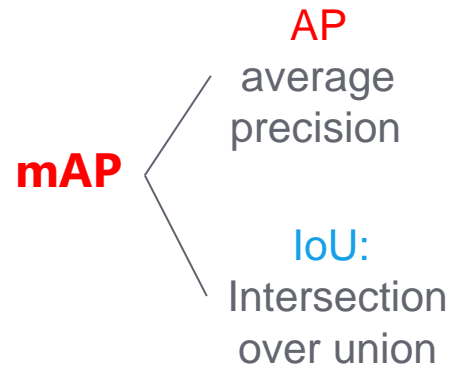


4 m



Evaluation Metrics

1° Step: License Plate Detection



$$AP = \frac{\text{Correct predictions}}{\text{Total n° License Plates}}$$

How many plates were correctly located

$$IoU = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$




How well the prediction is overlapping to the LP



2° Step: Text Recognition: **Accuracy** = $\frac{\text{Correct license plates}}{\text{Total n° License Plates}}$

A license plate is correct if all the characters are correctly recognized
One mistake = wrong recognition



1 B A 3 1 0 2
Correct!



1 B A 2 1 0 2
Wrong!

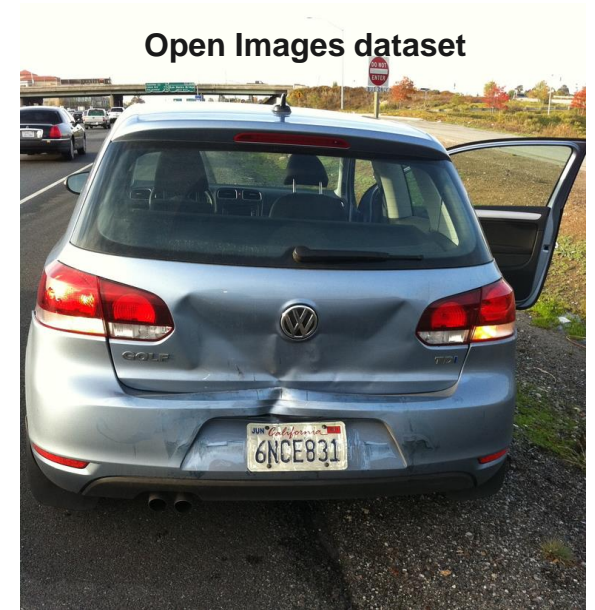
Design Flow on GAP

Step 1: License Plate Detection

SSD - Object Detection

Design flow of the license plate detector:

1. **Dataset search:** open source Google OpenImages
7k high-res images
2. Find the right **architecture** (accuracy/memory tradeoff)
Resnet: 270Mb (too big)
MobilenetV1-V2: 22,7Mb ; 18,5Mb (still too big)
SSDlite MobilenetV2: 13,3 Mb
3. Optimization: **quantization aware training**
float32bit → int8bit (4x memory save)
13,3 Mb → 3,2Mb
4. **Validation** on images acquired on gap
Not a problem transferring learning from hi-res images to Himax QVGA



My validation dataset (Himax pictures)



Training Experiments

Bigger images = better accuracy:

- SSD 512x512 images: 45.4%
- SSD 320x240 images (QVGA): 42.0% ↓

We are forced to use QVGA format (320x240).
But a camera with higher resolution improves accuracy.

Shrinking down the model's size:

Architectures	SSD MobilenetV1 (22Mb)	39.6%
	SSD MobilenetV2 (18Mb) ↑	42.0% ↑
	SSDlite MobilenetV2 (13.3Mb)	39.8% ↓

We want to fit in 8 MB

Failed approaches:

- Grayscale training: 38.3% ↓
- Frozen backbone training: 17.3% ↓

Validation images are grayscale.
Transform training images to grayscale too

Start from a pretrained model and just
finetune last layer of the network

Optimization: quantization aware training


- SSD SSDliteV2 (13.3Mb) 39.8%
- SSD SSDliteV2 8bit quant (3.2Mb) ↑ 38.9% ↓

Final model

Final Detection Accuracy

Model's final mAP score: **38.9 %**

- SSDlite MobilenetV2 architecture
- 3.2Mb
- 8 bit quantized
- Training on QVGA resolution images (320x240)



Result Comparison:

- Nice accuracy, same algorithm performs 22% on COCO.
- But here we are detecting a single class (easier problem)

Reference benchmark : standard COCO dataset

Algorithm	Architecture	Size	mAP on COCO
YOLOv2	Darknet 19	(60Mb)	21%
YOLOv3	Darknet 53	(190Mb)	33%
SSD	Mobilenet v1	(73Mb)	18%
SSD	Mobilenet v2	(138Mb)	22%
SSD	SSDlitev2	(48Mb)	22%
SSD	Resnet 50	(350Mb)	35%
SSD	Resnet 101	(610Mb)	40%

Detection distance

After 2 meters the characters are not readable anymore. But the algorithm still works.

0.3 m



1 m



2 m



3 m



4 m



Design Flow on GAP

Step 2: Character Recognition

LPRNet – License Plate Recognition Network

Design flow for character recognition:

1. Dataset search:

Synthetic Chinese License Plates (270k images)

CCPD: real Chinese License Plates (200k images)

Czech Low-Quality License Plates (185k images)

2. Architecture manipulation:

original size: 6.8Mb

replaced unsupported operations: 9.6Mb

Depth reduction fully connected layers : 4 Mb

3. Optimization: quantization aware training

float32bit → int8bit (4x memory save)

4 Mb → 1.0Mb

4. Validation on images acquired on gap

Himax QVGA grayscale images

Training Datasets for char recognition



GAP Validation image



Accuracy on benchmark datasets

Synthetic Chinese License Plates: 97.7 %



CCPD: real Chinese License Plates: 99.1%



Czech Low-Quality License Plates: 99.2%



Accuracy after quantization and depth reduction

Training and testing on Synthetic chinese license plates

Original network (6.8Mb):	98.1 %	optimization ↓
Replace unsupported operations (9.6Mb):	97.5% ↓	
Depth reduction fully connected layers (4 Mb):	97.6% —	
Quantization aware training (1.0Mb):	97.7% —	

There is **no accuracy drop** while **optimizing**: brought down the model size from 9.6 to 1 Mb

Maximum Recognition Distance

QVGA (320x240) low quality images: from **2 meters** the characters are barely readable.

The algorithm is reliable up to 1m of distance.

0.3 m
Accuracy: 100%

1 m
Accuracy: 100%

2 m
Accuracy: 36.4%
Accur-1 : 54.5%

3 m
Accuracy: 0%



<Shanghai>A708K1



<Shanghai>A708K1



<Shanghai>A708K1



<Hebei>ATQ8KV



<Shanghai>AFM3883



<Shanghai>AFM3883



<Zhejiang>AFM3883



<Hebei>PE3



<Shandong>Q3X5U3



<Shandong>Q3X5U3




<Shandong>Q3X5U3



<Shandong>Q6

Good resistance to tilting and blur

- Tilting Effects



<Hebei>A6171A



<Hubei>HJ9108



<Anhui>HZ381W

- Blurred images



1BA3102



1BA3102

Final Pipeline

Unification of the 2 steps

Final Unified Pipeline

Unification of the 2 steps:

1. SSD: License plate detection (3.2Mb)
2. LPRNet: Text Recognition (1Mb)

Final LP model size : 3.2 Mb + 1Mb = **4.2 Mb OK!**

Gap memory constrain: **8Mb**

So it can fit within L3 memory of GAP

Input Image



1 Step: SSD
for plate detection



2 Step: LPRNet
for text recognition



<Henan>K0F755

Final Pipeline Speed: does it run in real-time?

Inference estimation

		Worst case 1 MAC/Cycle		best case 8 MAC/Cycle
1° Step (SSD):	540M MACs	→	3.6 sec at 150MHz	→ 0.45 sec at 150MHz
2° Step (LPRNet):	172M MACs	→	1.8 sec at 150MHz	→ 0.14 sec at 150MHz
Steps Combined:	712M MACs	→	5.4 sec at 150MHz	→ 0.49 sec at 150MHz

Real-Time applications: ~ 1 FPS

Our Speed:

0.49 sec < Inference Speed < 5.4sec

0.2 FPS < Inference Speed < 2 FPS

Conclusion

Conclusion

2 Steps approach for License Plate OCR:

1. License Plate Detection: SSD (Single Shot Detector)
2. License Plate Text Recognition: LPRNet

Works on Himax QVGA grayscale images:

- Detects License plates up to **4m** distance.
- Recognize Text up to **1m** distance.

mAP score = 38.9%

Accuracy = 97.7 - 99.2%

Final LP model size : 3.2 Mb + 1Mb = 4.2 Mb

Fits in the L3 RAM memory of gap: < 8Mb

Inference speed (expected on GAP): 0.49 sec < time < 5.4sec

Could run in Real-Time: ~ 1 FPS

Input Image



License Plate Detection



License Plate Text Recognition



<Henan>K0F755

Thank you!

Questions?

Thanks for helping:

- Manuele Rusci
- Marco Fariselli
- Francesco Paci
- Ahmad Bijar

